




# Deep Visual Learning

---


**Jayanta Mukhopadhyay**  
**Dept. of Computer Science and Engg.**

# Deep learning

- Learning using a “deep” neural network





Classical ANN:  
Only a few  
hidden layers.



Deep architecture: Many hidden layers.

# Deep learning for solving vision problems

- Object recognition
- Object Localization
- Semantic Segmentation
- Video summarization
- Tracking objects
- ...





# Deep architecture: Why so late in application?

---

- Concepts introduced in 80's.
- Basic principles remain the same.
- Two major reasons.
  - Availability of large scale annotated data.
    - Penetration of internet and smart phones.
    - Wide spread of social networking.
    - Online shopping, etc.
  - Advancement of computing power.
    - High throughput GPU computing.

# Classical Image Classification



hand-crafted  
feature  
extractor

Classifier  
Algorithm

output

Tiger?

Cat?

Lion?

- Edges
- SIFT/SURF key Point
- HOG Regional Features
- Motion Features, etc.

- Bayesian
- LDA
- SVM
- KNN

# Classification Challenges

- Very tedious and costly to develop hand-crafted features to handle various challenges.



View Point variation



Deformation



Occlusion



Intraclass Variation



Illumination



Clutter



Instances

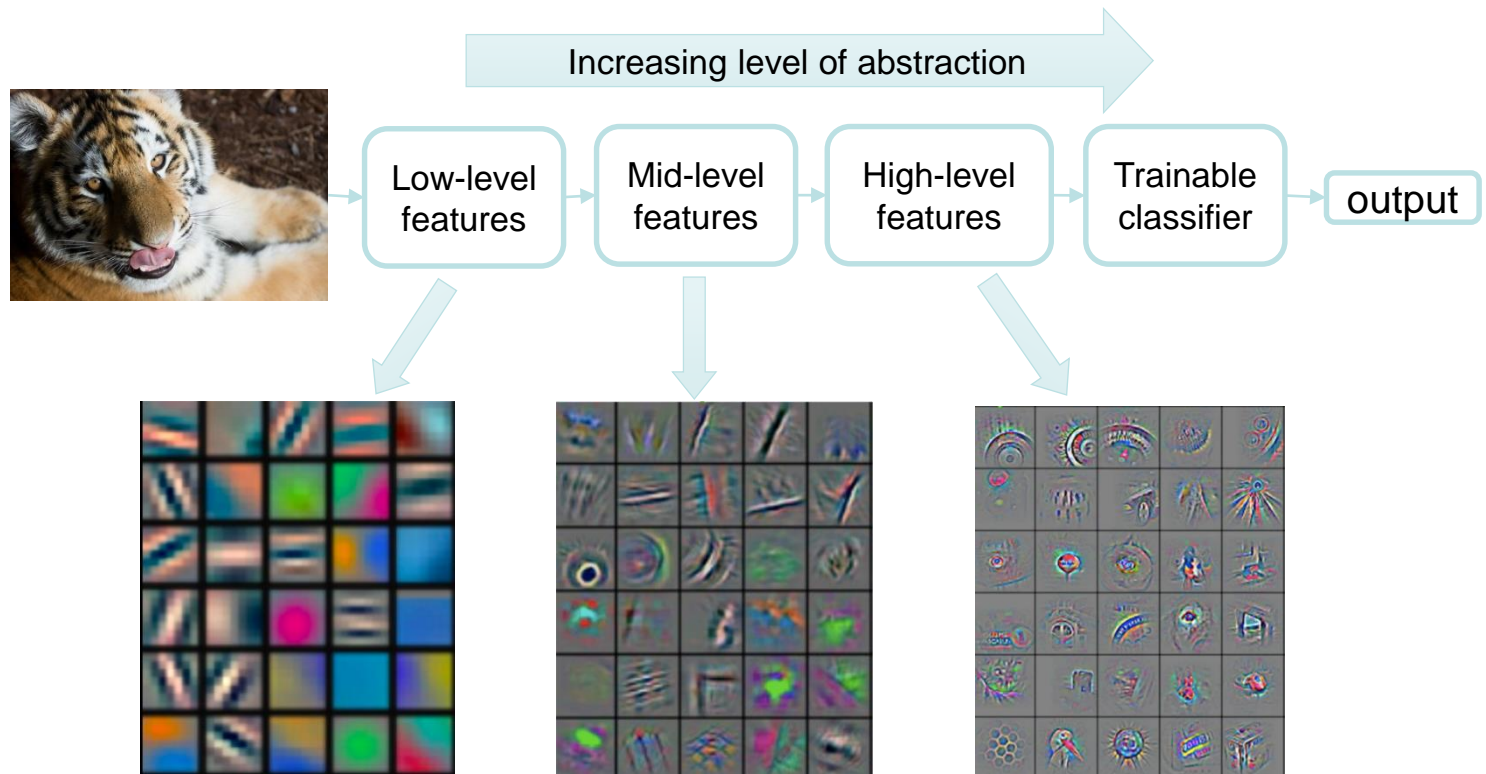


Scale

Highly dependent on one application, and not transferable easily to other applications.

# Classification through deep learning

Learning filters of feature extraction and also classifier.



\*Feature visualization of convolutional net trained on ImageNet (Zeiler and Fergus, 2013)



# Supervised Learning

---

**Data:**  $(x, y)$  where  $x$  is data,  $y$  is label

**Goal:** Learn a function  $f$  to map  $x \rightarrow y$

**Examples:** Classification, Regression,  
Object detection, Semantic  
Segmentation, Image Captioning, etc.





# Supervised Learning

---

Data Driven Approach to learn the model in three steps:

## Step 1: Define Model

$$\hat{y} = f(x, w)$$

Predicted output  
(image label)

Model  
structure

Input data  
(Image pixels)

Model weights

*Learn a parametric function  $f$  composed by weight parameters  $w$  to classify Image  $x$  as class label  $y$ .*

# Supervised Learning



---

**Step 2:** Collect data.

$$\{(x_i, y_i)\}_{i=1}^N$$

Training  
input

True  
output

# Supervised Learning

**Step 3:** Learn the model.

Total Loss = Data Loss + Regularization Loss

$$w^* = \arg \min_w \frac{1}{N} \sum_{i=1}^N \ell(\underbrace{f(x_i, w)}_{\text{Predicted output}}, y_i) + R(w)$$

**Learned weights** (points to  $w^*$ )

**Minimize average loss over training set** (points to  $\frac{1}{N} \sum_{i=1}^N$ )

**Loss function:** Measures “badness” of prediction (points to  $\ell$ )

**Regularizer:** Penalizes complex models (points to  $R(w)$ )



# Loss

---

- A **loss function** tells how good our current classifier is.
- **Data loss:** Model predictions should match training data
  - **Softmax Loss** (Multinomial Logistic Regression):

$$L_i = -\log\left(\frac{e^{s_{y_i}}}{\sum_j e^{s_j}}\right)$$

$$L_i = -\log(P(Y = y_i | X = x_i))$$

# Cross-entropy Loss

- Another form of softmax loss.

- **2-class entropy:**

- $-(y \log(p) + (1-y) \log(1-p)); \quad p: \text{Prob. } (y=1/\text{o})$

- **Multiclass:**

$$-\sum_{c=1}^M y_{\text{o},c} \log(p_{\text{o},c})$$

Estimated Prob. of **o** belonging to **c**

Binary indicator (1 if **o** belongs to **c**, else 0).

True Prob. of **o** belonging to **c**

**More general:**

$$-\sum_{c=1}^M q_{\text{o},c} \log(p_{\text{o},c})$$



# Regularization Loss

---

*i*

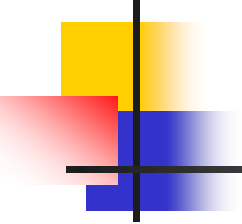
■ **Regularization Loss:** Model should be “simple”, so it works on test data as “ $W$ ” is not unique with just data loss.

■  $L_2$  Regularization (Weight Decay)  $R(W) = \sum_k \sum_l W_{k,l}^2$

■  $L_1$  Regularization  $R(W) = \sum_k \sum_l |W_{k,l}|$

■ Elastic net ( $L_1 + L_2$ )  $R(W) = \sum_k \sum_l \beta W_{k,l}^2 + |W_{k,l}|$

# How to find best weights $w^*$ ?


$$w^* = \arg \min_w \frac{1}{N} \sum_{i=1}^N \ell(f(x_i, w), y_i) + R(w)$$

$\underbrace{\hspace{15em}}_{g(w)}$

$$= \arg \min_w g(w)$$

- Gradient descent
  - Back propagation algorithm



# Gradient Descent

## How to update weights?

Initialize  $w$  randomly

While true:

    Compute gradient  $\nabla g(w)$  at current point

    Move downhill a little bit:  $w = w - \alpha \nabla g(w)$

updating the weights at each  
iteration

**Learning rate:** How big  
each step should be





# Back Propagation

---

- **Forward pass:**

- Run graph “forward” to compute loss

- **Backward pass:**


- Run graph “backward” to compute gradients with respect to loss

- Efficient to compute gradients for big, complex models.

# Learning filters for feature extraction

- Correlation with a mask or kernel


$w_1$	$w_2$	$w_3$
$w_4$	$w_c$	$w_5$
$w_6$	$w_7$	$w_8$



$$g(x, y) = w_1 f(x-1, y+1) + w_2 f(x, y+1) + w_3 f(x+1, y+1) + w_4 f(x-1, y) + w_c f(x, y) + w_5 f(x+1, y) + w_6 f(x-1, y-1) + w_7 f(x, y-1) + w_8 f(x+1, y-1)$$


# Convolution in neural architecture

- Output of a neuron: weighted sum of inputs
  - weights defined by a kernel
  - Sparse connectivity
- Shared weights for every node
  - Sufficient to describe the model by a kernel



# Convolution Layer

32x32x3 image



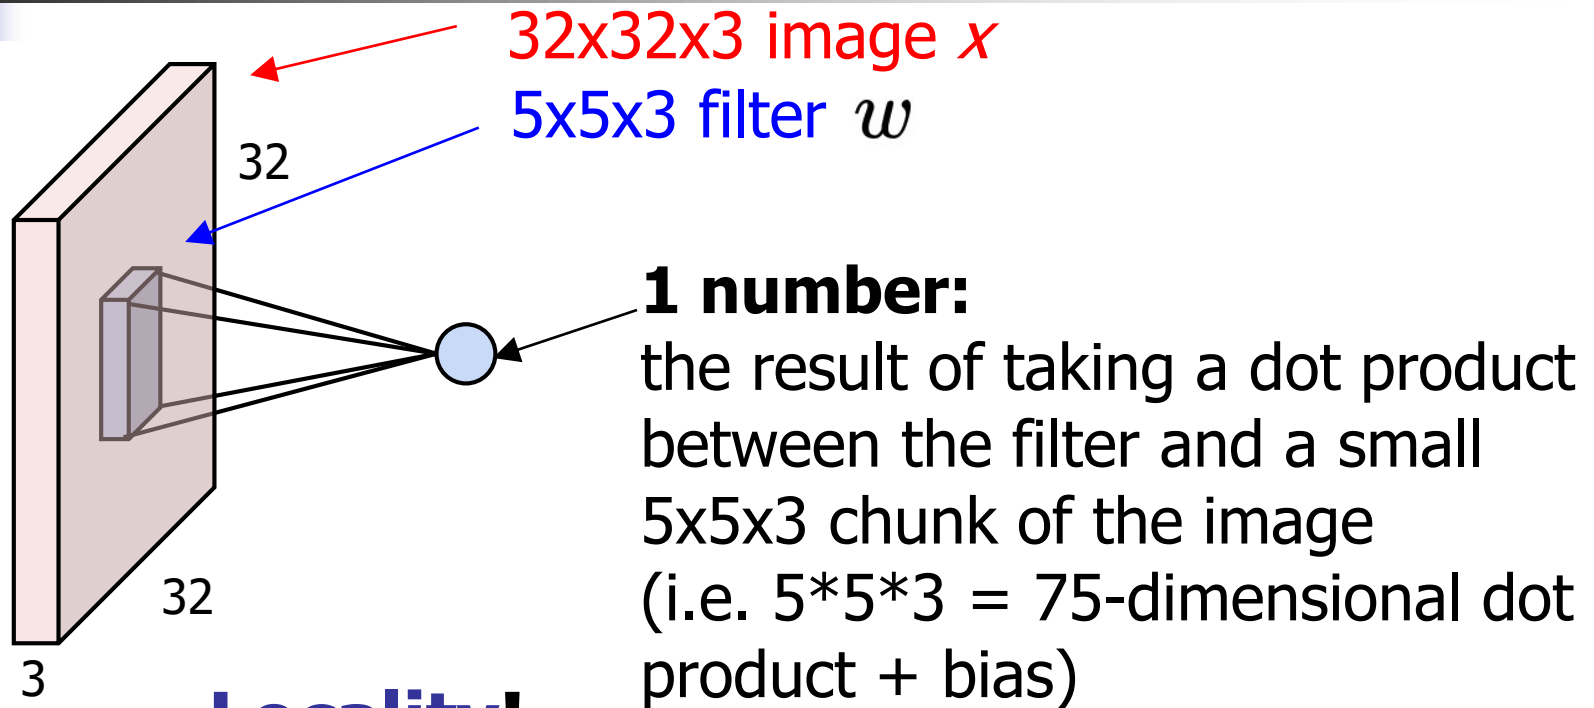
Filters always extend the full depth of the input volume

5x5x3 filter (kernel)



**Convolve** the filter with the image  
i.e. "slide over the image spatially,  
computing dot products".

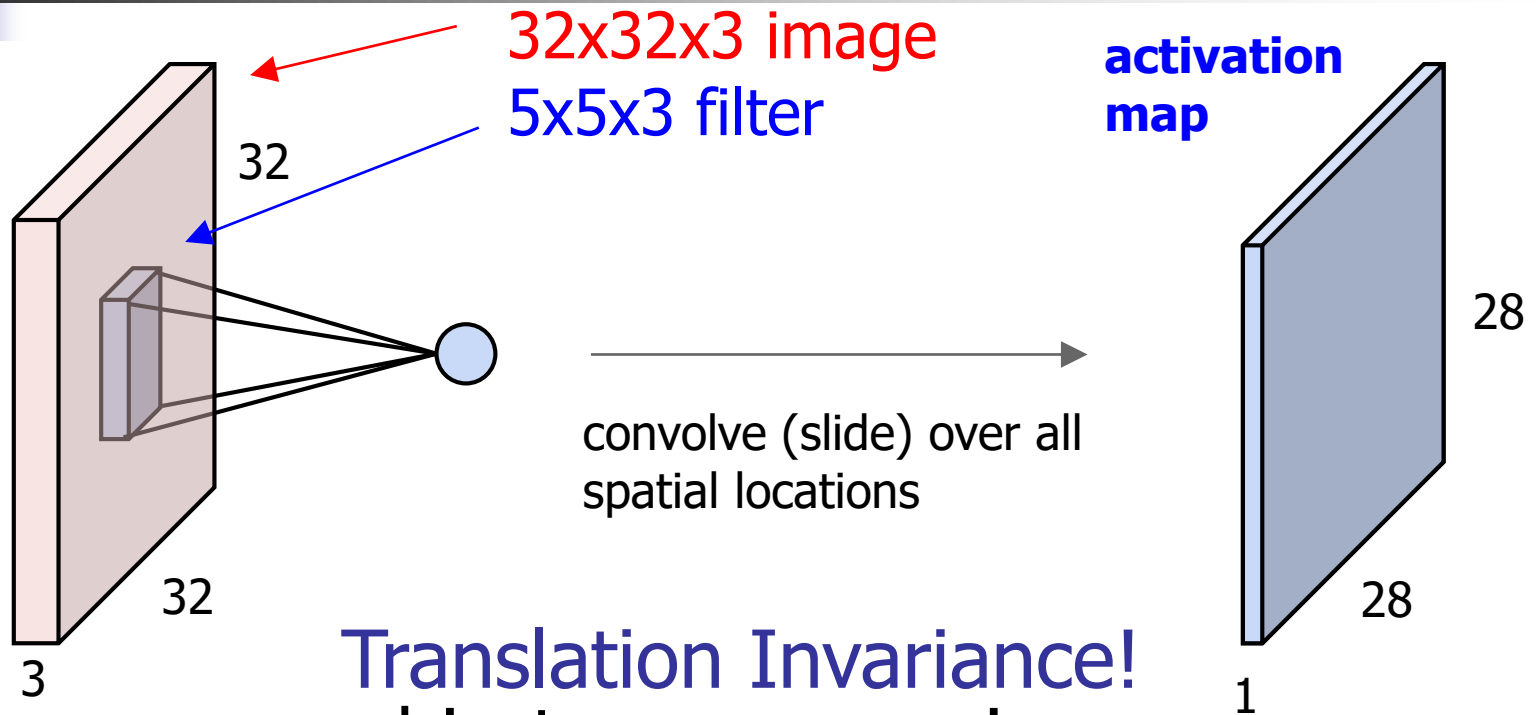
# Convolution Layer



$$w^T x + b$$


Objects tend to have a local  
spatial support.

# Convolution Layer



**Translation Invariance!**  
object appearance is  
independent of location  
**Weight sharing!**


# Convolution Layer



Consider a second, **green** filter.

# Convolution Layer (CONV)

**activation maps**



For example, if we had 6  $5 \times 5 \times 3$  filters, we'll get 6 separate activation maps:

We stack these up to get a "new image" of size  $28 \times 28 \times 6$ !





# Features of CONV

---

- **Locality:**
  - objects tend to have a local spatial support
- **Translation invariance:**
  - object appearance is independent of location
- **Weight sharing**
  - units connected to different locations have the same weights
  - equivalently, each unit is applied to all locations
  - weights of filters are invariant.
- Each unit output of filter is connected to a local rectangular area in the input.
  - – Receptive Field



# Non-Linear Layer


---

- Increase the nonlinearity of the entire architecture without affecting the receptive fields of the convolution layer.
  - Commonly used in CNN is **ReLU**.


# Non Linearity: Activation Functions

## - A few examples


**Sigmoid**  $\sigma(x) = 1/(1 + e^{-x})$




**tanh(x)**




**ReLU** (Rectified Linear Unit)



**Leaky ReLU**



# Convolutional Neural Networks (CNN)



A CNN is a sequence of convolution layers and nonlinearities.




# Parameters involved in convolution layer

---

- Input Volume size  $W_1 \times H_1 \times D_1$
- No. of filters  $K$  with size  $F_w \times F_h \times D_1$  convolved with stride  $(S_w, S_h)$ .
- Input zero padded by  $(P_w, P_h)$  on both sides.
- **Output volume size  $W_2 \times H_2 \times D_2$ ?**
  - $W_2 = (W_1 - F_w + 2P_w)/S_w + 1$
  - $H_2 = (H_1 - F_h + 2P_h)/S_h + 1$
  - $D_2 = K$
- **Parameters ?**
  - $(F_w * F_h * D_1) * K$  weights +  $K$  biases
- **d-th depth slice of output is the result of convolution of d-th filter over the padded input volume with a stride, then offset by d-th bias**




# Pooling Layer (POOL)




- To progressively reduce the spatial size of the representation.
  - to reduce the amount of parameters and computation in the network.
  - to control overfitting.
- Pooling partitions the input image into a set of non-overlapping rectangles.
- For each such sub-region, outputs an aggregated value of the features in that region.
  - Maximum value (Max pooling)
  - Average value (Average pooling)
- Operates over each activation map independently

# Pooling Layer (POOL)


224x112x64



Pool



112x56x64



1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

Maxpool with 2x2  
filter and stride 2

6	8
3	4

Single depth slice



# Parameters involved in pooling

---

- Input Volume size  $W_1 \times H_1 \times D_1$
- Pool size  $F_w \times F_h$  with stride  $(S_w, S_h)$ .
- Output volume size  $W_2 \times H_2 \times D_2$ ?
  - $W_2 = (W_1 - F_w) / S + 1$
  - $H_2 = (H_1 - F_h) / S + 1$
  - $D_2 = D_1$
- Parameters ?
  - 0!
- **Uncommon to use zero-padding in Pooling layers.**





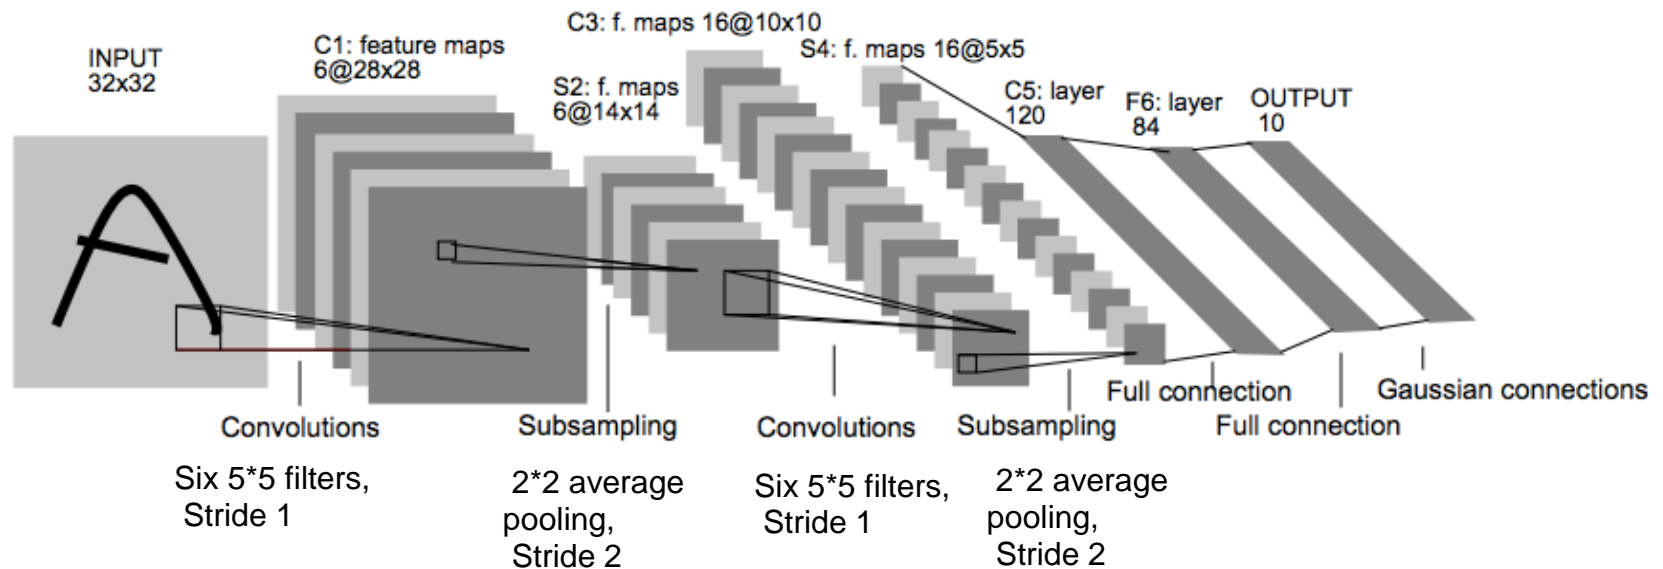
# Fully Connected Layer (FC)

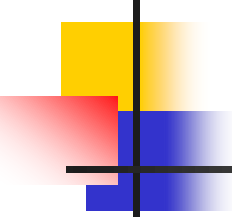
---

- Contains neurons that connect to the entire input volume
  - as in ordinary Neural Networks.
- Input volume to FC layer can also be treated as Deep Features.
- If the FC layer is a classifier, the input to FC can also be treated as feature vector representation for the sample.

# LeNet: A typical example

- I/P → CONV → POOL → CONV → POOL → FC → FC → O/P
  - Number of parameters: 60k
  - Number of floating point operations per inference: 341k
  - Sigmoid used for non-linearity.





# Efficient computation with smaller kernels

---

- Successive filtering with smaller sized kernels covers equivalent receptive field area of a larger size.
  - Stack of three 3x3 conv (stride 1) layers has **same effective field** as one 7x7 conv layer
  - Deeper with more non-linearity
    - Usually deeper the model better the accuracy
    - Till it overfits!
  - Fewer parameters:
    - $3 \cdot (3^2 C^2)$  vs.  $(7^2 C^2)$  for C channels
- Used in VGG network
  - no. of layers: 13 / 16 / 19 with only 3x3 kernels

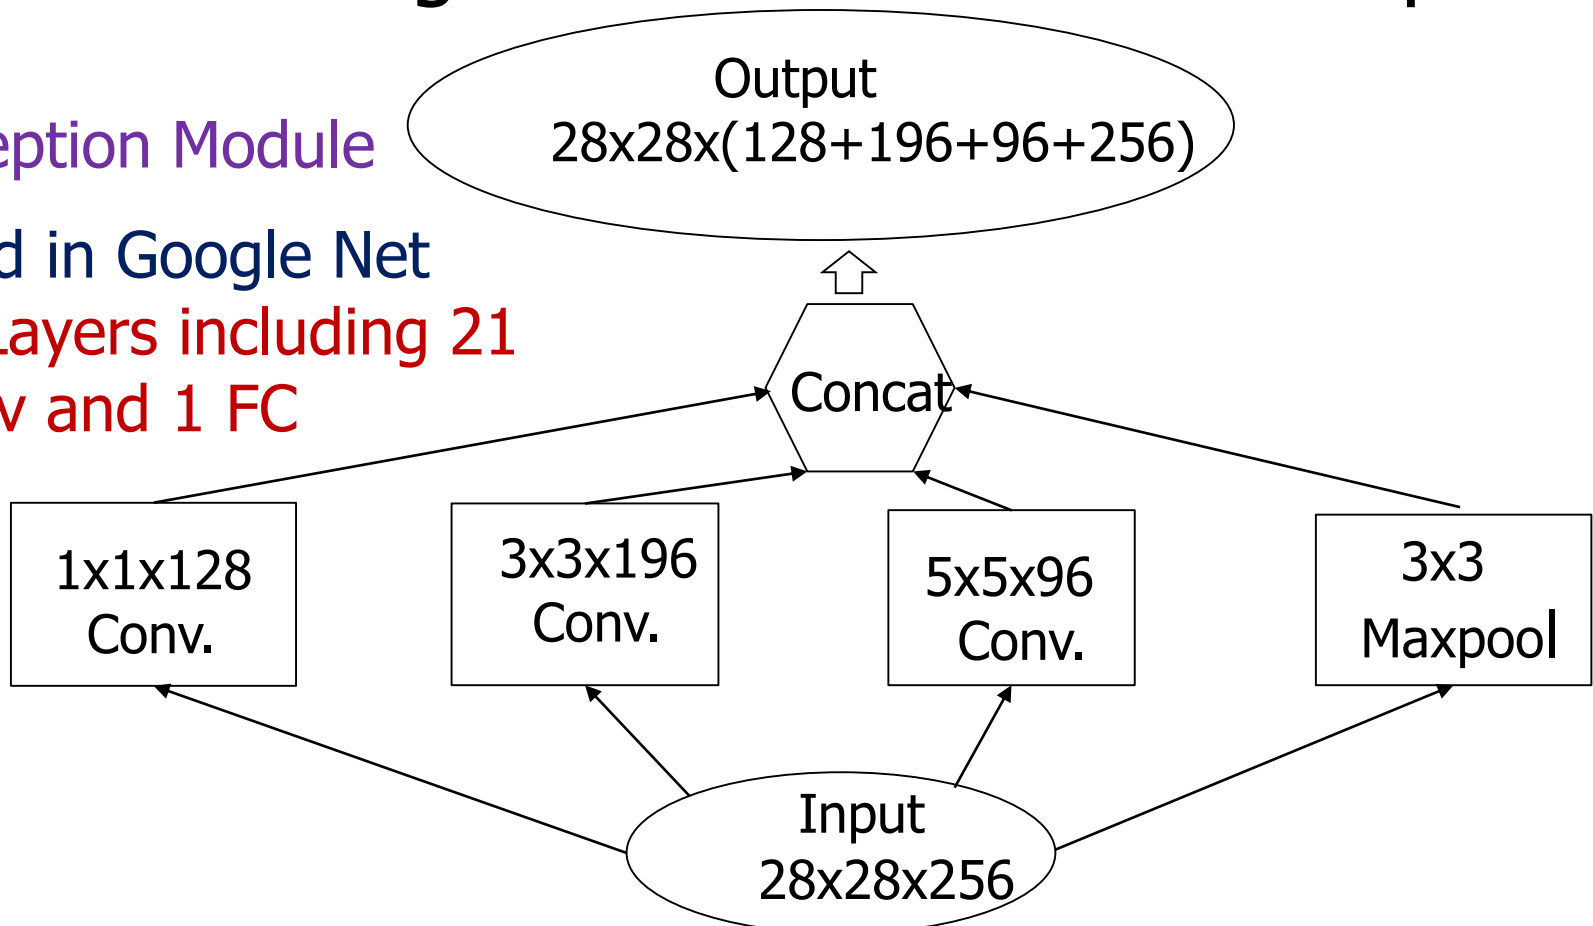
# Handling scale in feature representation

- Concatenating multiscale feature descriptor

Inception Module

Used in Google Net


57 Layers including 21  
Conv and 1 FC



# Vanishing gradient problem

- Gradient becomes zero (vanishes) at deeper layers!
- Learn residual mapping!

Used in ResNet  
No. of layers:  
34/50/101/152






# Batch Normalization

---

- Normalizes input activation map to a layer by considering its distribution over a batch of training samples.
  - Each dimension of the input feature map individually normalized
  - To make Gaussian activation maps.
- Advantages
  - Improves gradient flow through the network.
  - Allows higher learning rates.
  - Reduces the strong dependence on initialization.
  - Acts as a form of regularization.
- Usually inserted after FC / CONV layers, and before non-linearity.



**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

**Algorithm 1:** Batch Normalizing Transform, applied to activation  $x$  over a mini-batch.



# Drop out

---

- Randomly dropping out nodes of network (at hidden / visible layers) during training.
  - Temporarily removing it from the network, along with all its incoming and outgoing connections.
  - To regulate overfitting, more effective for smaller dataset.
  - Simulates learning sparse representation in hidden layers.
- Implementation
  - Retain output of a node with a probability  $p$ .
    - Typically within  $[0.5, 1]$  at hidden layers and  $[0.8, 1]$  in visible layers.






# Learning weights with drop out

---

- Weights become larger due to drop out.
  - Needs to be scaled at the end training.
  - A simple heuristic.
    - Outgoing weights of a unit retained with probability  $p$  during training, multiplied by  $p$  at test time.
  - Scaling may be carried out during training time at each weight update.
    - No need to rescale weight for the test network.

# Depthwise Separable Convolutions

- Suppose, we have  $D_F \times D_F \times M$  input feature map,  $D_F \times D_F \times N$  output feature map and  $D_k \times D_k$  spatial sized conventional convolution filters.



- What is the computational cost for such a convolution operation?

$$— D_k \cdot D_k \cdot M \cdot D_F \cdot D_F \cdot N$$


- What is the number of parameters?

$$— D_k \cdot D_k \cdot M \cdot N$$

# Depthwise Separable Convolutions

❑ Now, think of  $M$  filters which are  $D_K \times D_K$  (not  $D_K \times D_K \times M$ ) and think each  $M$  of these filters are operated separately on  $M$  channels of input of spatial size  $D_F \times D_F$

❑ Number of parameters?  $D_K \cdot D_K \cdot M$




❑ What is the computational cost for such a convolution operation?  
 $D_K \cdot D_K \cdot D_F \cdot D_F \cdot M$

❑ This operation is known as **Depthwise Convolution** operation.

# Depthwise Separable Convolutions

- ❑ What is the output shape now?  $D_F \times D_F \times M$
- ❑ Where did the  $N$  (output channels) go?
  - not there as depthwise convolution operates only on input channels.
- ❑ Now think about  $1 \times 1$  traditional convolution on  $D_F \times D_F \times M$  featuremap to get  $D_F \times D_F \times N$  output.
- ❑ What is the computation cost?  $1 \cdot 1 \cdot M \cdot D_F \cdot D_F \cdot N = D_F \cdot D_F \cdot M \cdot N$



Used in MobileNet-V1  
Alternate layers of  
Conv and D-S Conv

What is the number of parameters?  $1 \cdot 1 \cdot M \cdot N$

- ❑ This operation is called  $1 \times 1$  pointwise convolution

Courtesy: Ankita Chatterjee

# MobileNet-V1

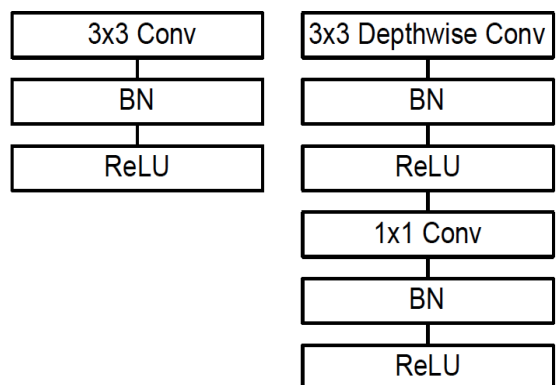


Figure 3. Left: Standard convolutional layer with batchnorm and ReLU. Right: Depthwise Separable convolutions with Depthwise and Pointwise layers followed by batchnorm and ReLU.

Image taken from: *MobileNet Paper*

Table 1. MobileNet Body Architecture

Type / Stride	Filter Shape	Input Size
Conv / s2	$3 \times 3 \times 3 \times 32$	$224 \times 224 \times 3$
Conv dw / s1	$3 \times 3 \times 32 \text{ dw}$	$112 \times 112 \times 32$
Conv / s1	$1 \times 1 \times 32 \times 64$	$112 \times 112 \times 32$
Conv dw / s2	$3 \times 3 \times 64 \text{ dw}$	$112 \times 112 \times 64$
Conv / s1	$1 \times 1 \times 64 \times 128$	$56 \times 56 \times 64$
Conv dw / s1	$3 \times 3 \times 128 \text{ dw}$	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 128$	$56 \times 56 \times 128$
Conv dw / s2	$3 \times 3 \times 128 \text{ dw}$	$56 \times 56 \times 128$
Conv / s1	$1 \times 1 \times 128 \times 256$	$28 \times 28 \times 128$
Conv dw / s1	$3 \times 3 \times 256 \text{ dw}$	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 256$	$28 \times 28 \times 256$
Conv dw / s2	$3 \times 3 \times 256 \text{ dw}$	$28 \times 28 \times 256$
Conv / s1	$1 \times 1 \times 256 \times 512$	$14 \times 14 \times 256$
$5 \times$	Conv dw / s1	$3 \times 3 \times 512 \text{ dw}$
	Conv / s1	$1 \times 1 \times 512 \times 512$
	Conv dw / s2	$3 \times 3 \times 512 \text{ dw}$
	Conv / s1	$1 \times 1 \times 512 \times 1024$
	Conv dw / s2	$3 \times 3 \times 1024 \text{ dw}$
	Conv / s1	$1 \times 1 \times 1024 \times 1024$
	Avg Pool / s1	Pool $7 \times 7$
	FC / s1	$1024 \times 1000$
	Softmax / s1	Classifier
		$1 \times 1 \times 1000$

Image taken from: *MobileNet Paper*

Courtesy: Ankita Chatterjee

# Width and Resolution Multiplier

- ❑ Used as parameters for scaling the model architectures.
- ❑ Width multiplier  $\alpha \in (0, 1]$  to thin a network uniformly at each layer
  - The number of input channels (from  $M$ ):  $\alpha M$
  - The number of output channels (from  $N$ ):  $\alpha N$
  - Computational cost?  $D_k \cdot D_k \cdot \alpha M \cdot D_F \cdot D_F + D_F \cdot D_F \cdot \alpha M \cdot \alpha N$
  - Reduces roughly by  $\alpha^2$
- ❑ Resolution multiplier  $\rho \in (0, 1]$  to reduce the image resolution and the internal representation of every layer by by this factor
- ❑ With width multiplier  $\alpha$  and resolution multiplier  $\rho$ , the computational cost?  $D_k \cdot D_k \cdot \alpha M \cdot \rho D_F \cdot \rho D_F + \rho D_F \cdot \rho D_F \cdot \alpha M \cdot \alpha N$ 
  - Another reduction by  $\rho^2$



# MobileNet-V1

Table 4. Depthwise Separable vs Full Convolution MobileNet

Model	ImageNet	Million	Million
	Accuracy	Mult-Adds	Parameters
Conv MobileNet	71.7%	4866	29.3
MobileNet	70.6%	569	4.2

Image taken from: *MobileNet Paper*

Table 6. MobileNet Width Multiplier

Width Multiplier	ImageNet	Million	Million
	Accuracy	Mult-Adds	Parameters
1.0 MobileNet-224	70.6%	569	4.2
0.75 MobileNet-224	68.4%	325	2.6
0.5 MobileNet-224	63.7%	149	1.3
0.25 MobileNet-224	50.6%	41	0.5

Table 7. MobileNet Resolution

Resolution	ImageNet	Million	Million
	Accuracy	Mult-Adds	Parameters
1.0 MobileNet-224	70.6%	569	4.2
1.0 MobileNet-192	69.1%	418	4.2
1.0 MobileNet-160	67.2%	290	4.2
1.0 MobileNet-128	64.4%	186	4.2

Image taken from: *MobileNet Paper*



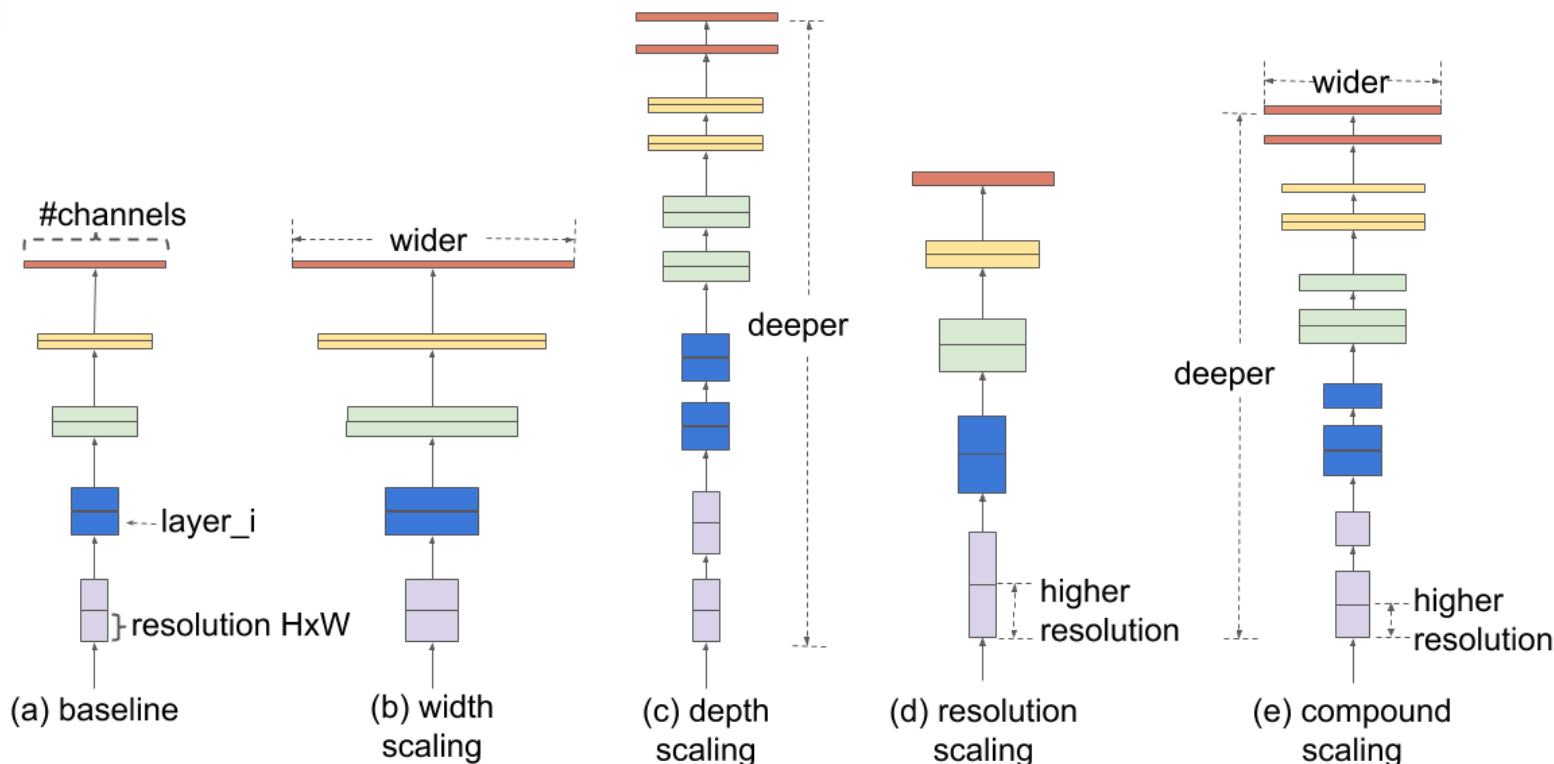
# EfficientNets: Compound Model Scaling

---

- ❑ By balancing all dimensions of the network—width, depth, and image resolution—against the available resources to get the best overall performance.
- ❑ To perform a grid search to find the relationship between different scaling dimensions of the baseline network under a fixed resource constraint
- ❑ Apply those coefficients to scale up the baseline network to the desired target model size or computational budget



# Compound Model Scaling



Comparison of different scaling methods. Unlike conventional scaling methods (b)-(d) that arbitrarily scale a single dimension of the network, compound scaling method uniformly scales up all dimensions in a principled way.

Image taken from: *EfficientNet* Paper

Courtesy: Ankita Chatterjee

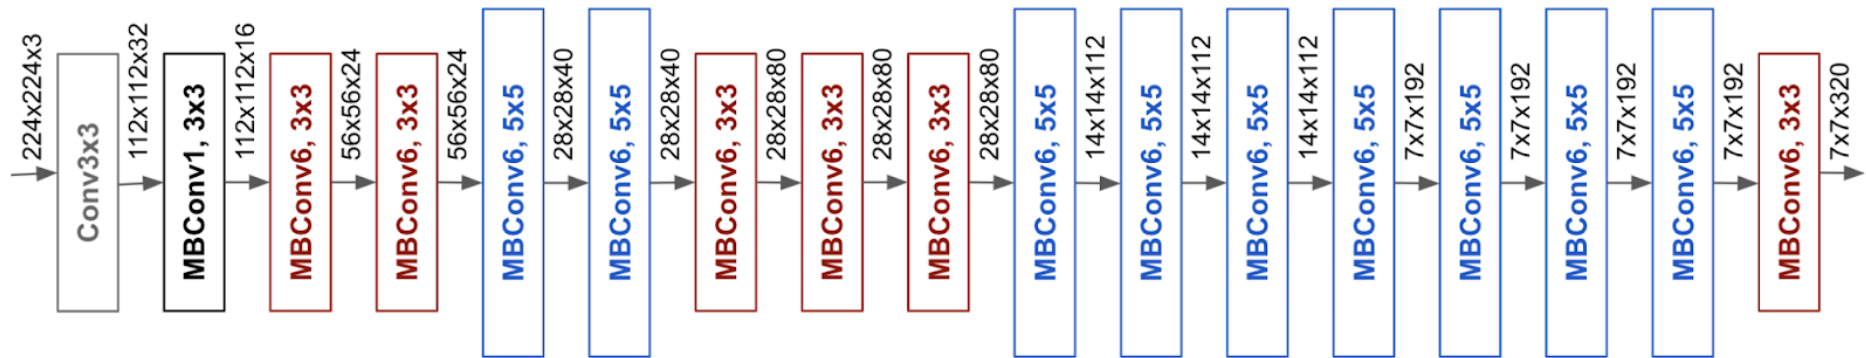


# EfficientNet Architecture

---


- ❑ The effectiveness of model scaling also relies heavily on the baseline network
- ❑ A new baseline network is developed by performing a neural architecture search using the AutoML MNAS framework, which optimizes both accuracy and efficiency (FLOPS).
- ❑ The resulting architecture uses mobile inverted bottleneck convolution (MBConv), similar to MobileNetV2 and MnasNet, but is slightly larger

# EfficientNet Architecture



The architecture for baseline network EfficientNet-B0 is simple and clean, making it easier to scale and generalize.  
Image taken from: *EfficientNet* Paper

# EfficientNet Architecture



## Model Size vs. Accuracy

EfficientNet-B0:  
the baseline network  
developed by AutoML MNAS,

Efficient-B1 to B7:  
obtained by scaling up the  
baseline network.

In particular, EfficientNet-B7  
achieves new state-of-the-art  
84.4% top-1 / 97.1% top-5  
accuracy, while being 8.4x  
smaller than the best existing  
CNN.

Courtesy: Ankita Chatterjee

# EfficientNet performance on other baseline architectures

Model	FLOPS	Top-1 Acc.
Baseline MobileNetV1 (Howard et al., 2017)	0.6B	70.6%
Scale MobileNetV1 by width ( $w=2$ )	2.2B	74.2%
Scale MobileNetV1 by resolution ( $r=2$ )	2.2B	72.7%
<b>compound scale (<math>d=1.4, w=1.2, r=1.3</math>)</b>	<b>2.3B</b>	<b>75.6%</b>
Baseline MobileNetV2 (Sandler et al., 2018)	0.3B	72.0%
Scale MobileNetV2 by depth ( $d=4$ )	1.2B	76.8%
Scale MobileNetV2 by width ( $w=2$ )	1.1B	76.4%
Scale MobileNetV2 by resolution ( $r=2$ )	1.2B	74.8%
<b>MobileNetV2 compound scale</b>	<b>1.3B</b>	<b>77.4%</b>
Baseline ResNet-50 (He et al., 2016)	4.1B	76.0%
Scale ResNet-50 by depth ( $d=4$ )	16.2B	78.1%
Scale ResNet-50 by width ( $w=2$ )	14.7B	77.7%
Scale ResNet-50 by resolution ( $r=2$ )	16.4B	77.5%
<b>ResNet-50 compound scale</b>	<b>16.7B</b>	<b>78.8%</b>

Scaling Up MobileNets and ResNet.

# EfficientNet performance vs. other models

Model	Top-1 Acc.	Top-5 Acc.	#Params	Ratio-to-EfficientNet	#FLOPs	Ratio-to-EfficientNet
<b>EfficientNet-B0</b>	<b>77.1%</b>	<b>93.3%</b>	<b>5.3M</b>	<b>1x</b>	<b>0.39B</b>	<b>1x</b>
ResNet-50 (He et al., 2016)	76.0%	93.0%	26M	4.9x	4.1B	11x
DenseNet-169 (Huang et al., 2017)	76.2%	93.2%	14M	2.6x	3.5B	8.9x
<b>EfficientNet-B1</b>	<b>79.1%</b>	<b>94.4%</b>	<b>7.8M</b>	<b>1x</b>	<b>0.70B</b>	<b>1x</b>
ResNet-152 (He et al., 2016)	77.8%	93.8%	60M	7.6x	11B	16x
DenseNet-264 (Huang et al., 2017)	77.9%	93.9%	34M	4.3x	6.0B	8.6x
Inception-v3 (Szegedy et al., 2016)	78.8%	94.4%	24M	3.0x	5.7B	8.1x
Xception (Chollet, 2017)	79.0%	94.5%	23M	3.0x	8.4B	12x
<b>EfficientNet-B2</b>	<b>80.1%</b>	<b>94.9%</b>	<b>9.2M</b>	<b>1x</b>	<b>1.0B</b>	<b>1x</b>
Inception-v4 (Szegedy et al., 2017)	80.0%	95.0%	48M	5.2x	13B	13x
Inception-resnet-v2 (Szegedy et al., 2017)	80.1%	95.1%	56M	6.1x	13B	13x
<b>EfficientNet-B3</b>	<b>81.6%</b>	<b>95.7%</b>	<b>12M</b>	<b>1x</b>	<b>1.8B</b>	<b>1x</b>
ResNeXt-101 (Xie et al., 2017)	80.9%	95.6%	84M	7.0x	32B	18x
PolyNet (Zhang et al., 2017)	81.3%	95.8%	92M	7.7x	35B	19x
<b>EfficientNet-B4</b>	<b>82.9%</b>	<b>96.4%</b>	<b>19M</b>	<b>1x</b>	<b>4.2B</b>	<b>1x</b>
SENet (Hu et al., 2018)	82.7%	96.2%	146M	7.7x	42B	10x
NASNet-A (Zoph et al., 2018)	82.7%	96.2%	89M	4.7x	24B	5.7x
AmoebaNet-A (Real et al., 2019)	82.8%	96.1%	87M	4.6x	23B	5.5x
PNASNet (Liu et al., 2018)	82.9%	96.2%	86M	4.5x	23B	6.0x
<b>EfficientNet-B5</b>	<b>83.6%</b>	<b>96.7%</b>	<b>30M</b>	<b>1x</b>	<b>9.9B</b>	<b>1x</b>
AmoebaNet-C (Cubuk et al., 2019)	83.5%	96.5%	155M	5.2x	41B	4.1x
<b>EfficientNet-B6</b>	<b>84.0%</b>	<b>96.8%</b>	<b>43M</b>	<b>1x</b>	<b>19B</b>	<b>1x</b>
<b>EfficientNet-B7</b>	<b>84.3%</b>	<b>97.0%</b>	<b>66M</b>	<b>1x</b>	<b>37B</b>	<b>1x</b>
GPipe (Huang et al., 2018)	84.3%	97.0%	557M	8.4x	-	-

Image taken from: EfficientNet Paper

Courtesy: Ankita Chatterjee



# EfficientNet performance on other datasets using transfer learning

	Comparison to best public-available results						Comparison to best reported results					
	Model	Acc.	#Param	Our Model	Acc.	#Param(ratio)	Model	Acc.	#Param	Our Model	Acc.	#Param(ratio)
CIFAR-10	NASNet-A	98.0%	85M	EfficientNet-B0	98.1%	4M (21x)	<sup>†</sup> Gpipe	<b>99.0%</b>	556M	EfficientNet-B7	98.9%	64M (8.7x)
CIFAR-100	NASNet-A	87.5%	85M	EfficientNet-B0	88.1%	4M (21x)	Gpipe	91.3%	556M	EfficientNet-B7	<b>91.7%</b>	64M (8.7x)
Birdsnap	Inception-v4	81.8%	41M	EfficientNet-B5	82.0%	28M (1.5x)	GPipe	83.6%	556M	EfficientNet-B7	<b>84.3%</b>	64M (8.7x)
Stanford Cars	Inception-v4	93.4%	41M	EfficientNet-B3	93.6%	10M (4.1x)	<sup>‡</sup> DAT	<b>94.8%</b>	-	EfficientNet-B7	94.7%	-
Flowers	Inception-v4	98.5%	41M	EfficientNet-B5	98.5%	28M (1.5x)	DAT	97.7%	-	EfficientNet-B7	<b>98.8%</b>	-
FGVC Aircraft	Inception-v4	90.9%	41M	EfficientNet-B3	90.7%	10M (4.1x)	DAT	92.9%	-	EfficientNet-B7	<b>92.9%</b>	-
Oxford-IIIT Pets	ResNet-152	94.5%	58M	EfficientNet-B4	94.8%	17M (5.6x)	GPipe	<b>95.9%</b>	556M	EfficientNet-B6	95.4%	41M (14x)
Food-101	Inception-v4	90.8%	41M	EfficientNet-B4	91.5%	17M (2.4x)	GPipe	93.0%	556M	EfficientNet-B7	<b>93.0%</b>	64M (8.7x)
Geo-Mean	<b>(4.7x)</b>						<b>(9.6x)</b>					



# Training steps:

---

- Preprocessing of training dataset.
  - Normalize data.
  - Decorrelate data (Diagonal Covariance Matrix).
  - Whitening of data (Identity Covariance Matrix).
  - Subtract Mean.





# Training steps:

---

- Data augmentation.
  - Horizontal Flips
  - Random Crops on scaled input
  - Color jitter
  - Distortions
  - Transformations
- Weight initialization
- Train the network by update of the weight parameters.



# Few Training Tips

---

- Start with small regularization and find learning rate that makes the loss go down.
- Can overfit very small portion of the training data.
- Train first few epochs with few samples to initiate the hyper-parameters.
- If big gap between training accuracy and validation accuracy, then it is overfitting.
  - Try increase regularization.
- If no gap, then may increase model capacity.



# Transfer Learning


---

- No need of a lot of a data to train a CNN.
- Pre-trained models can be initialized for CNNs at the early stage of training.


# Transfer Learning

Lower learning rate when fine-tuning;  
1/10 of original LR good starting Point.

1. Train  
on Data  
Set #1




2. Train on  
smaller data set



Reinitialize  
this layer  
and train

Freeze  
these

3. Bigger  
dataset



With bigger  
dataset, train  
more layers

Train  
these  
layers

Freeze  
these

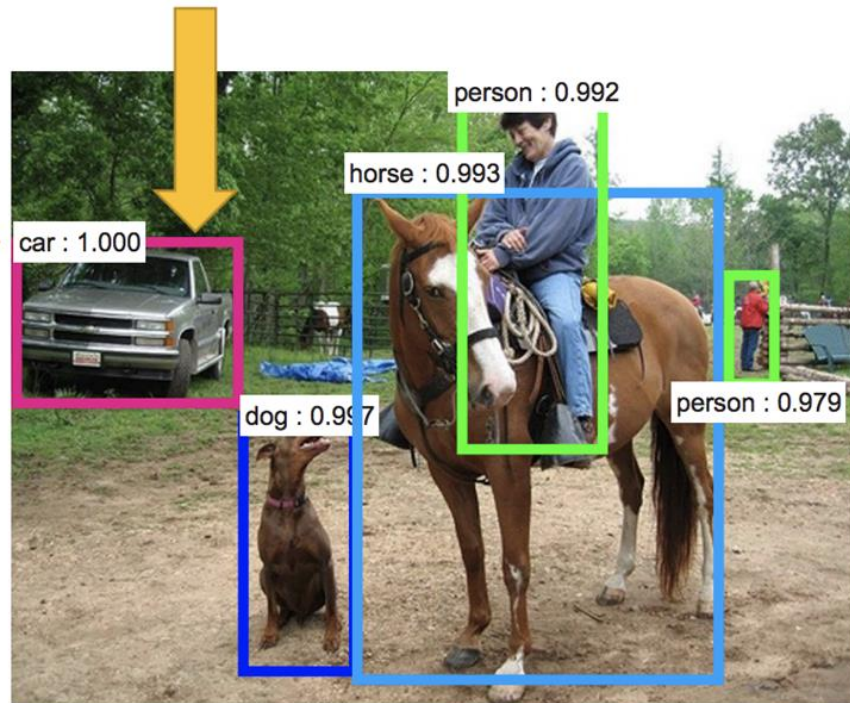
# Object Recognition and Localization

Recognition  
**What?**

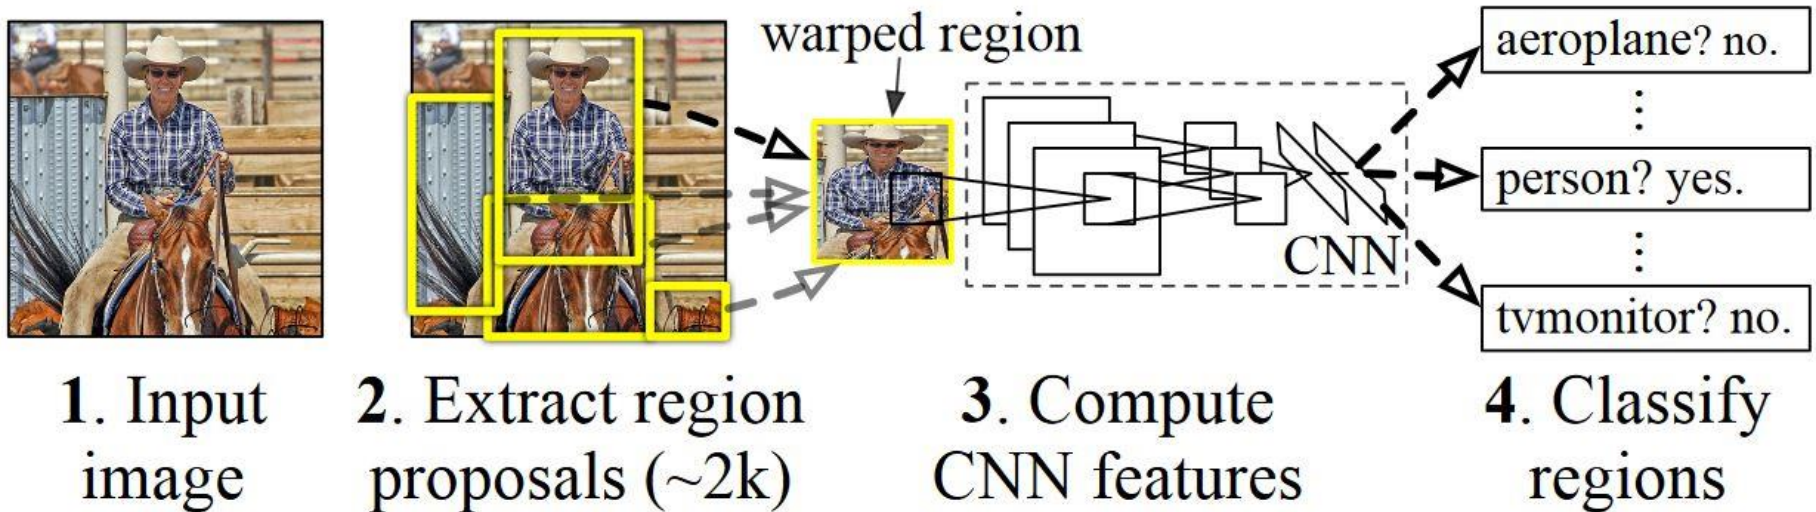
Localization  
**Where?**

Predicts

- (i) Coordinates of bounding boxes of objects.
- (ii) Class probs.



# Two stage processing: Localization and recognition








# Two stage processing: Object Recognition and Localization

---

	Region Proposal	Feature Extraction	Classification
Pre-CNN	Exhaustive	Hand Crafted	Linear
RCNN	Region Proposal	CNN	Linear SVM
Fast RCNN	Region Proposal	Deep	
Faster RCNN	Deep		

# Single stage processing:

## ■ YOLO (You Only Look Once)



1. Divide into  $S \times S$  non-overlapping grid.
2. For each grid, predict  $B$  boxes containing objects  $(x, y, w, h)$  with confidence score.
3. Get softmax probs. of object classes using a pre-trained network for the box with max. confidence with non-maximal suppression.

End to end training for both predicting bounding boxes and object classes.

Final output:  
 $S \times S \times (5B + C)$  tensor.





# YOLO architecture and loss function

---

- A series of convolution layers (coupled with max pooling) followed by a few fully connected layers.
  - Typically: For image of size 448x448
  - $S=7$ ,  $B=2$ ,  $C$  (no. of classes): 20
- Loss function:
  - Localization loss
  - Confidence loss
  - Classification loss
    - All of them use MSE.



# Semantic Segmentation

---

- Label each pixel in the image with a category label.
- Don't differentiate instances, only care about pixels.

## Instance Level Semantic Segmentation

- Even differentiate instances



# Semantic Segmentation

---

Building Blocks of CNNs:

- Convolution
- Down-Sampling
  - MaxPool, AvgPool, Strided Convolution ( $S > 1$ )
- Up-Sampling
  - UnPooling, Upconvolution


# Up- Sampling: Max Unpooling

1	2	<b>6</b>	3
3	<b>5</b>	2	1
1	2	2	1
<b>7</b>	3	4	<b>8</b>

MaxPool →

5	6
7	8

After few layers  
in Network



0	0	<b>b</b>	0
0	<b>a</b>	0	0
0	0	0	0
<b>c</b>	0	0	<b>d</b>

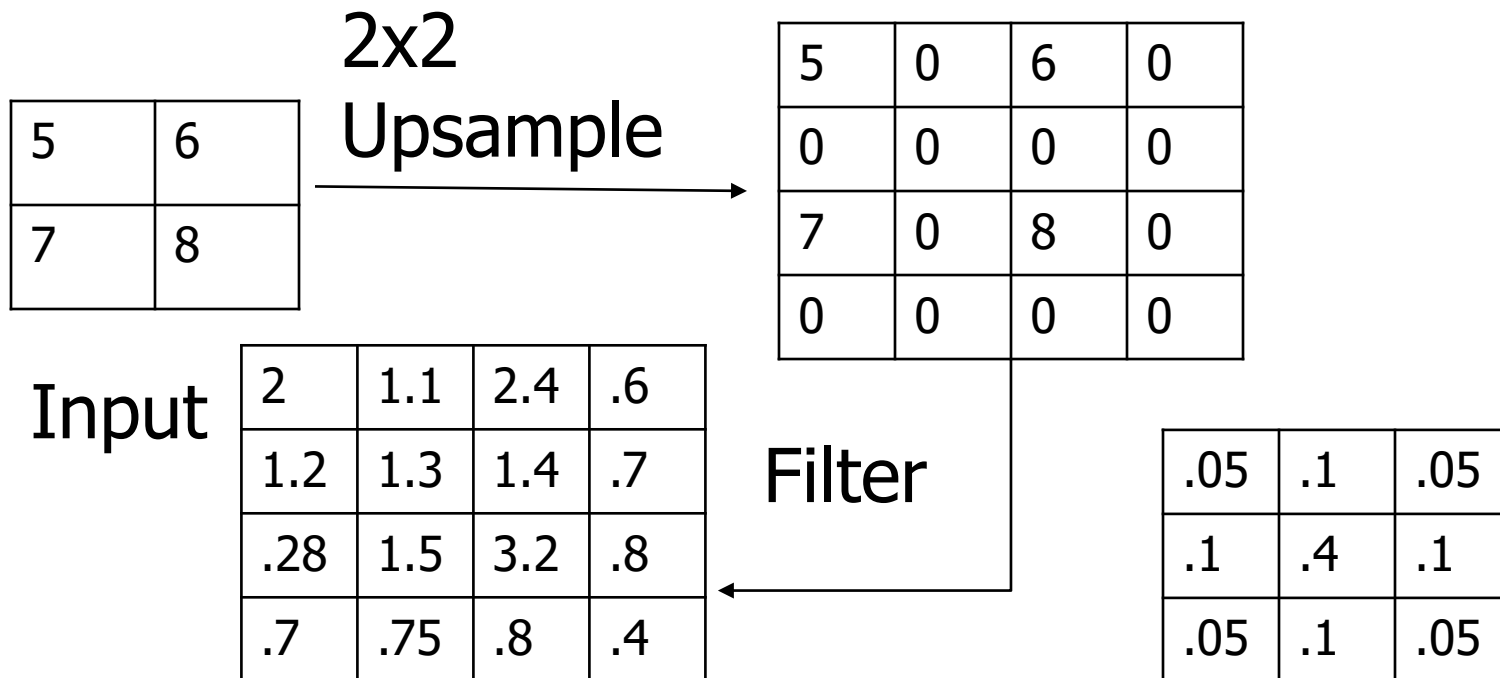
Max Unpool ←

a	b
c	d

0	0	<b>1</b>	0
0	<b>1</b>	0	0
0	0	0	0
<b>1</b>	0	0	<b>1</b>

Pooling Indices

# Upsampled convolution



# Fully Convolutional Network (FCN)

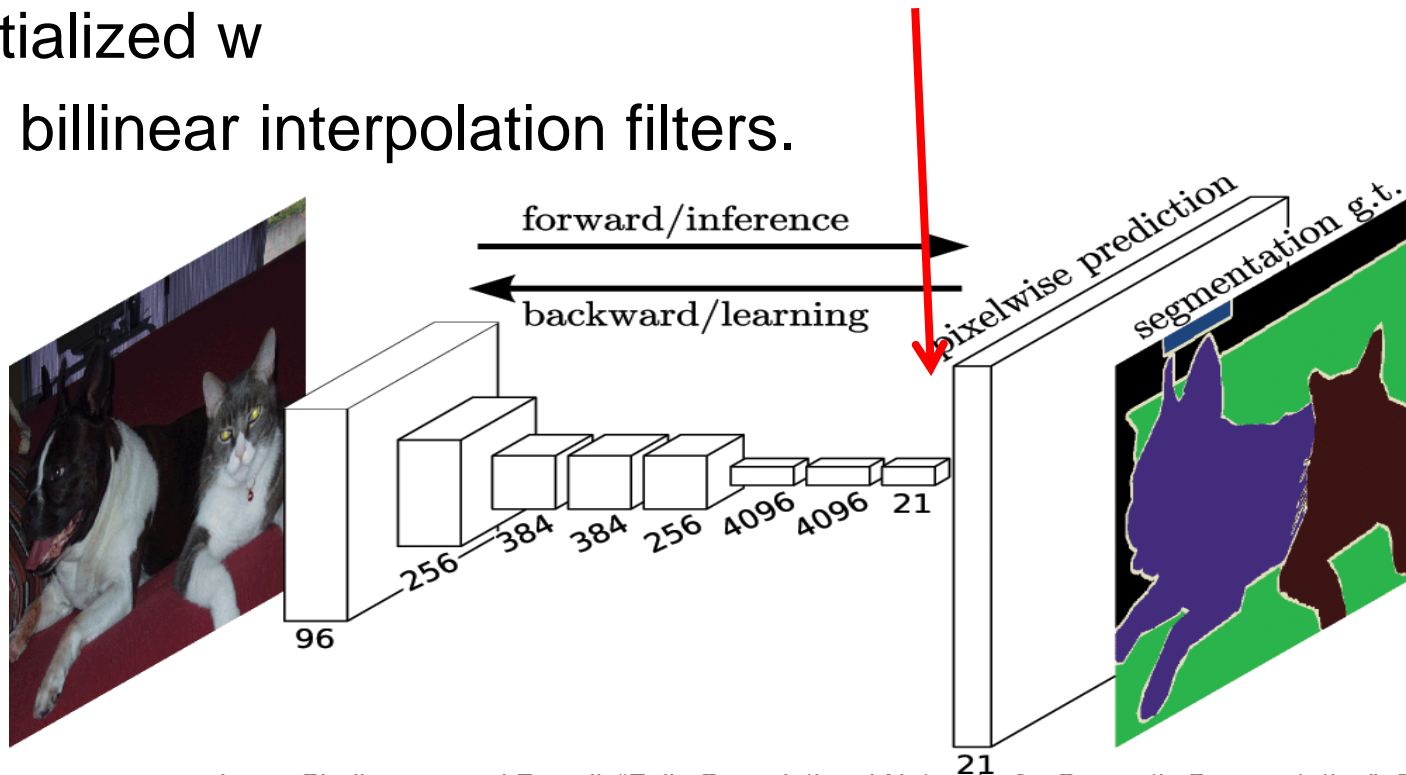


---

- No FC layer, only CNNs.
- Dense prediction
  - Downsampled output upsampled followed by 1x1 convolution for providing softmax score at every pixel.
  - No. of channel at the output layer=No. of classes.
  - Cross-entropy loss function.

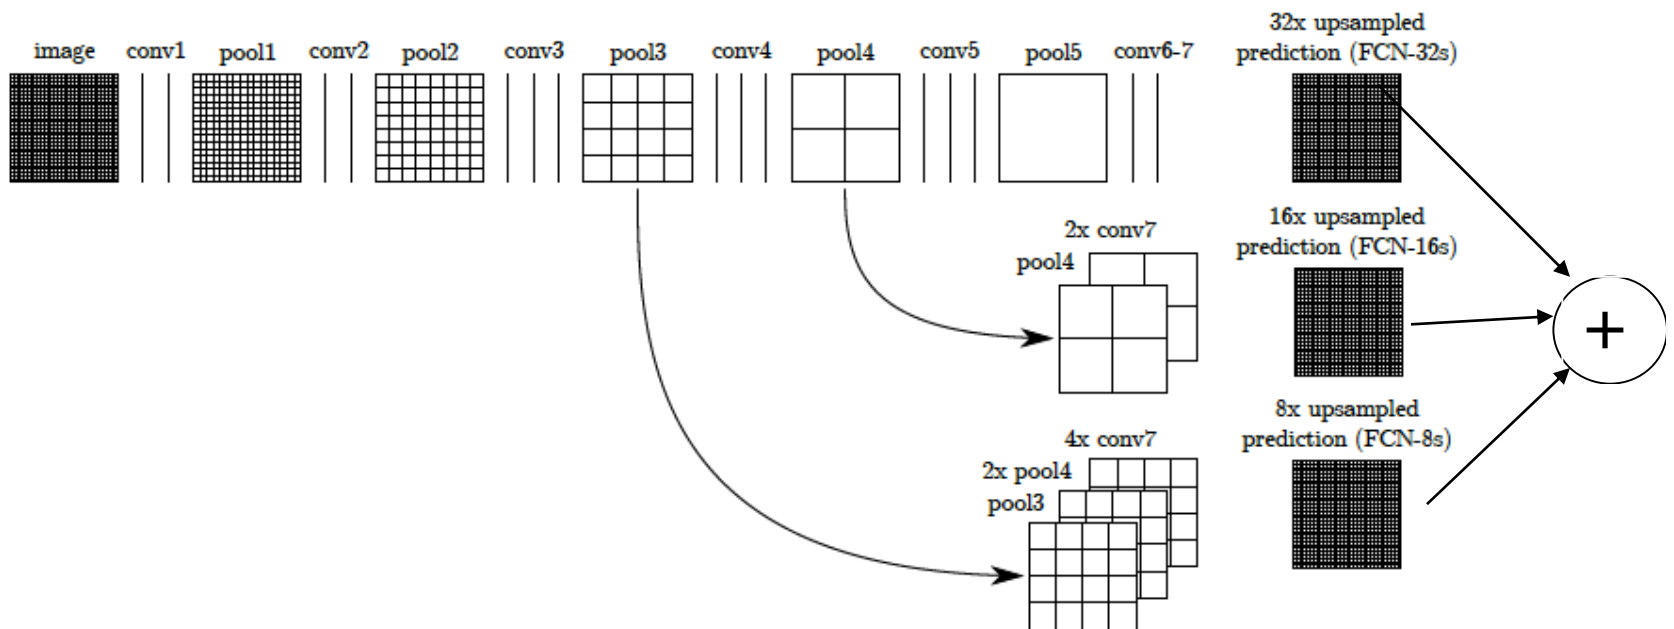
# Semantic Segmentation

- The upsampling of learned low resolution semantic feature maps is done using upconvolutions which are initialized with bilinear interpolation filters.



# Fusion of prediction

- Combines prediction of higher layer with lower layer (summing them).
- Performs at 3 stages.







# Segnet: Encoder-Decoder Network

---

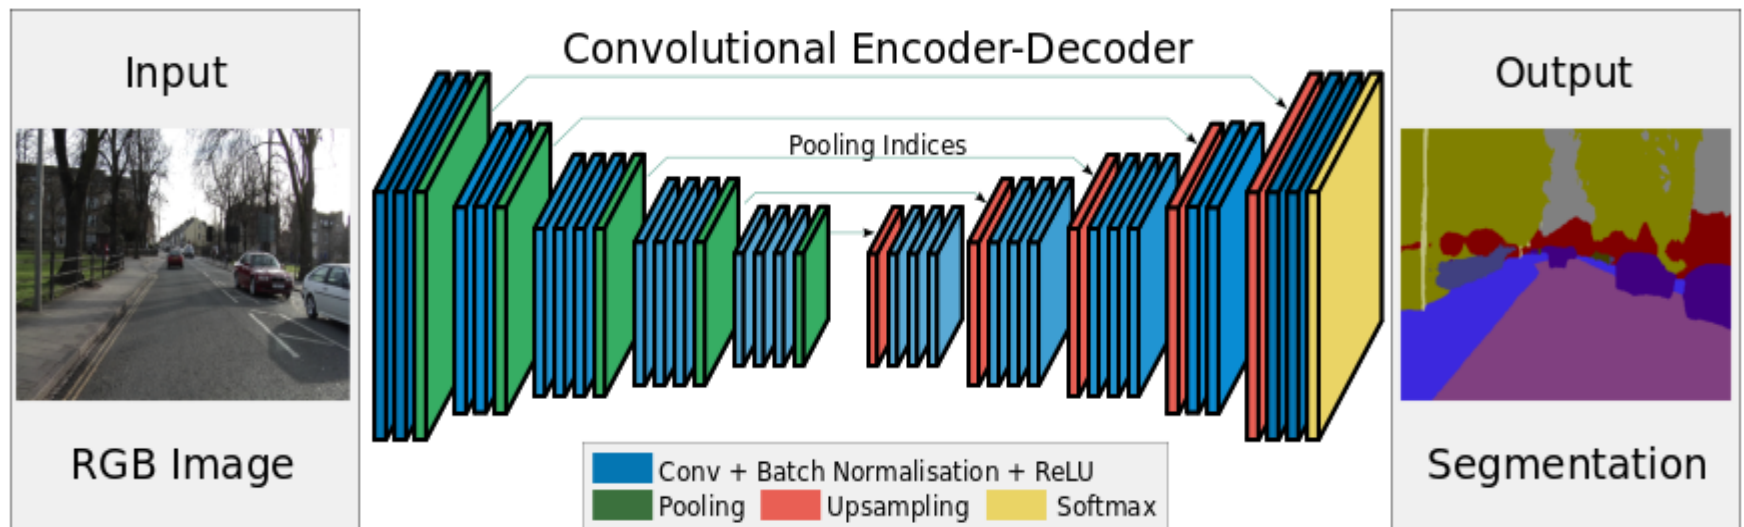
## Encoder

- Takes an input image and generates a high-dimensional feature vector
- Aggregate features at multiple levels

## Decoder

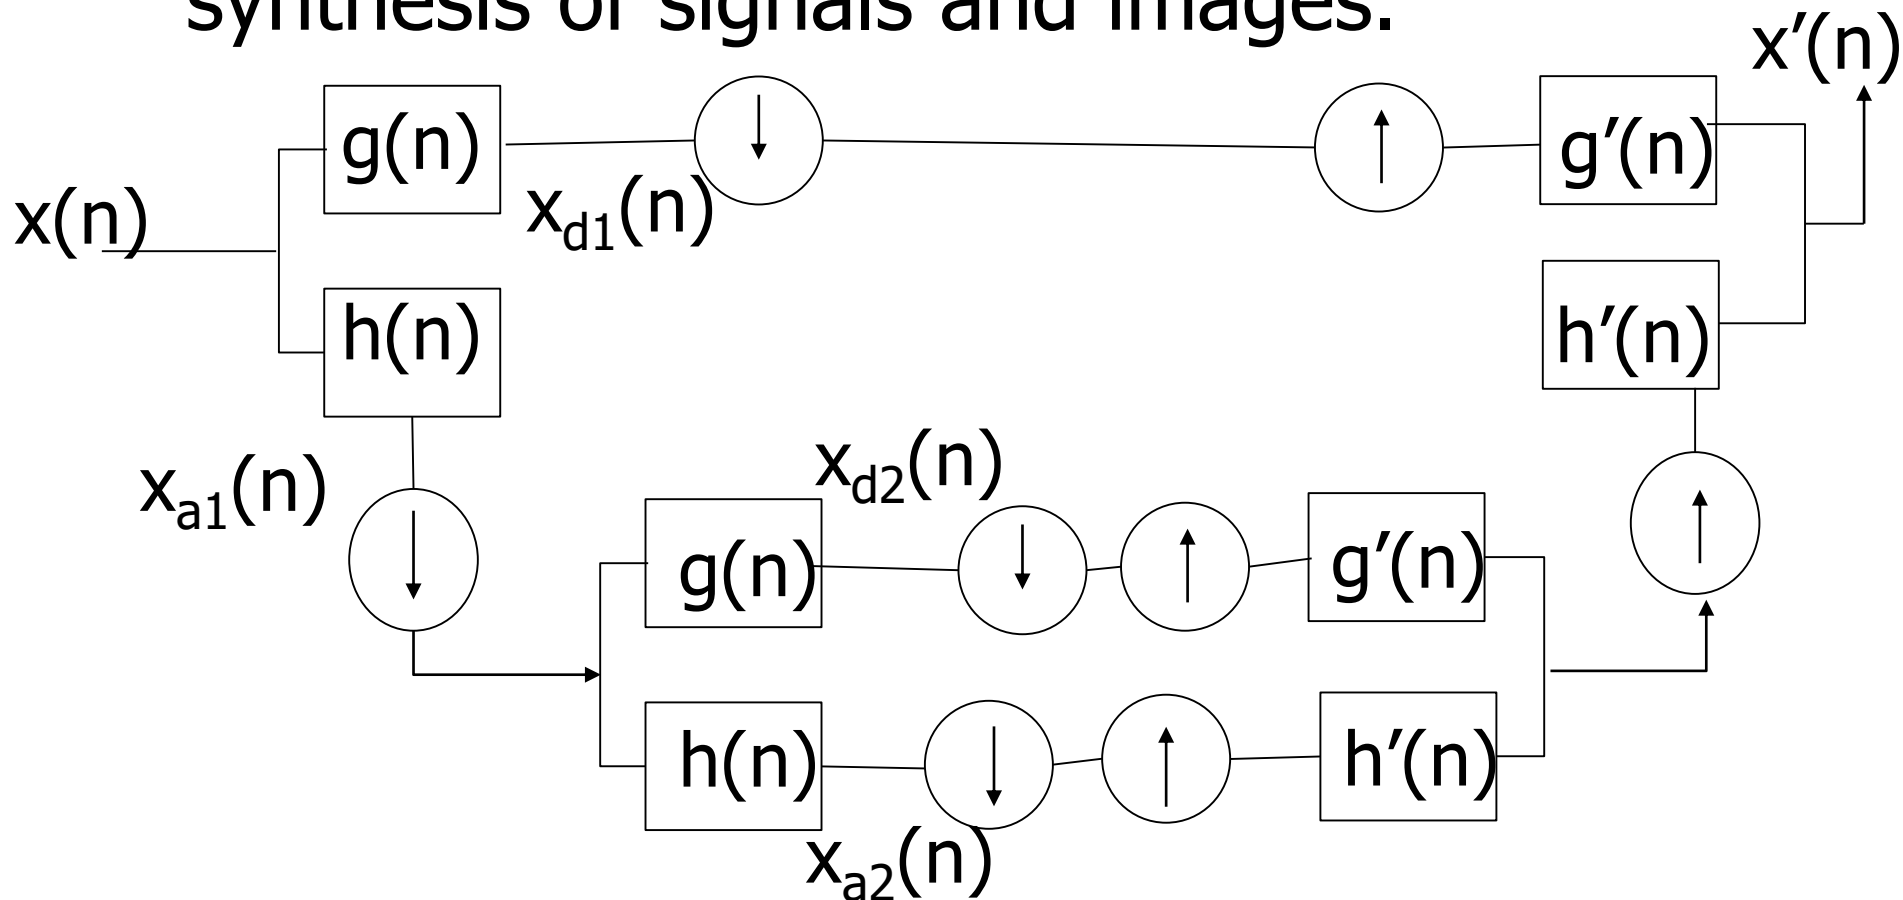
- Takes a high-dimensional feature vector and generates a semantic segmentation mask
- Decode features aggregated by encoder at multiple levels.
- Semantically project the discriminative features (lower resolution) learnt by the encoder onto the pixel space (higher resolution) to get a dense classification.

# Semantic Segmentation (Encoder-Decoder)




# U-Net

- Inspiration from wavelet analysis and synthesis of signals and images.



# U-Net architecture

- At the final layer a  $1 \times 1$  convolution used to map each feature vector to the desired number of classes.







# Results

Ground truth  
(Manual)

Yellow border

Color parts:


Segmented results



# Convolutional Encoder Decoder

Typical architecture

Self supervised  
Learning



Loss function: MSE

Auto encoder: Decoding the same image.



# Encoder - Decoder: Applications

---

- Denoising
- Segmentation
- Colorization
- Super-resolution
- ...

# Summary



---

- Deep architecture works in the same principle of artificial neural network.
  - A large number of hidden layers.
  - A large number of weights.
- Convolution Neural Network (CNN)
  - Learns filter weights.
    - Sharing of weights.
  - Two types of layers
    - Convolutional and Pooling Layers
  - Two stages
    - Feature extraction
    - Classification





# Summary

---

- Region proposal networks for simultaneous object localization and classification.
  - YOLO – end to end single staged network.
- Fully Convolutional Networks
  - Semantic Segmentation
  - U-Net
  - Encoder Decoder Network



Thank you!